# AI-Driven RF Design: An LLM Coding Agent Takes a Substrate Integrated Waveguide Filter from Specification to Fabrication

Federico Borello
Universidad Tecnológica Nacional — Facultad Regional Buenos Aires, Argentina
fborello.contact@gmail.com

***Abstract*—We report the design of a third-order, 2.87 GHz substrate integrated waveguide bandpass filter in which a large-language-model coding agent wrote and operated the entire toolchain: Chebyshev synthesis, parametric geometry, finite-element simulation with the open-source solver Palace, Bayesian optimization, and the export of fabrication files. One engineer set the specification, took the design decisions and ran the commercial-solver cross-checks. We report three results. First, a Gaussian-process optimizer applied to the assembled filter never reached a usable match, because the passband width it scores is discontinuous in the external coupling. Measuring smooth observables on small structures and inverting the textbook coupling relations converged in about three solves per coupling curve. Second, most of the 47 fixes in the project history corrected defects that produced plausible numbers without raising any error; we classify them and list the checks that now catch each class. Third, the fabricated design reaches a simulated centre frequency of 2.876 GHz, 182 MHz bandwidth, 1.70 dB insertion loss and -11.2 dB worst in-band return loss. It misses the -16 dB and 143.5 MHz targets, and a coupling-matrix fit attributes the gap mainly to end-cavity detuning. We also show that the coaxial connector is part of the tuning, and document a single-cavity dual-mode alternative as a negative result. The board is in manufacture; measured results will be added when it arrives.**

***Index Terms*—Substrate integrated waveguide, bandpass filter, large language models, AI agents, design automation, Bayesian optimization, finite-element method, coupling matrix.**

## Introduction

Machine learning has been part of microwave CAD for decades. Neural-network surrogate models [1], space mapping [2] and Bayesian optimization [3], [4] all replace some of the expensive electromagnetic (EM) solves with a cheaper learned or coarse model. Learned models are also used to extract coupling matrices from simulated responses, recently with physics-informed networks that need no labelled training set [5]. In each of these approaches the engineer still writes the geometry scripts, meshing, port setup, objective and export code, and interprets what the solver returns.

Recent large-language-model (LLM) agents work differently: they write and run code in a loop, observe the results and revise [6], [7]. Agent frameworks now exist for RFIC design [8] and RF amplifiers [9]. An agent has also operated CST, ADS and KiCad through their scripting interfaces to design a GNSS front end [10]. Empirical studies of bugs in agentic frameworks classify failures of the agent software itself [11]. What these reports leave largely unmeasured is how the *engineering* fails when an agent drives it: defects in the physical model that the agent writes, which matter more than the successes when one decides how much to trust the numbers.

This paper reports one complete design effort of that kind. The object is a classic structure, a three-cavity Chebyshev SIW filter [12], [13], [14]. It was chosen because its physics is textbook, so every agent-produced number can be checked against theory. An LLM coding agent (Claude Code [15]) wrote and ran all of the design software. The author acted as specifier, reviewer and decision-maker. The project history is preserved in git, and most claims below map to a commit or a stored result. Some process statistics, such as the screen-out rates of the Bayesian search, survive only in contemporaneous engineering notes rather than raw logs; we flag these where they appear.

The contributions are:

1. An account of the tool path: an FDTD solver (openEMS) that could not support match tuning, a finite-element solver (Palace) that could, and a GPU workstation run through remote agent sessions whose memory limits eventually pushed all production runs onto CPUs (Sec. III).
2. A measured explanation of why black-box Gaussian-process (GP) search over the assembled filter fails, and of why calibrate-and-invert succeeds (Sec. IV).
3. A taxonomy of the silent modelling defects the process produced and caught, with the guard that now blocks each class (Sec. V).
4. The fabricated design, with its deviation from specification diagnosed by coupling-matrix extraction. It is accompanied by a documented dual-mode negative result, an insertion-loss budget, a launch study and the manufacturing path (Secs. VI and VII).

## Filter specification and synthesis

The specification is f0 = 2.87 GHz, 5 % fractional bandwidth (FBW, 143.5 MHz), 0.1 dB Chebyshev ripple and $|S11| \leq -16$ dB in band. The laminate is Rogers RO4350B, 0.76 mm thick ($\varepsilon_r = 3.66$, $\tan\delta = 0.0037$ [16]), with 0.5 oz copper. The fence vias are 1.5 mm in diameter at a 3.0 mm pitch. Their p/d = 2 ratio sits at the edge of the leakage rules of [17], [18]. Smaller vias (0.8 mm at 1.5 mm pitch) were compared three times under three different objectives and never matched this choice. The connector is a Taoglas EMPCB.SMAFSTJ.B.HT edge-mount SMA [19].

The low-pass prototype gives g1 = g3 = 1.0316 and g2 = 1.1474 for N = 3 [20]. The bandpass mapping then gives the design targets:

$$k_{i,i+1} = \frac{\mathrm{FBW}}{\sqrt{g_i g_{i+1}}} = 0.0460, \qquad Q_e = \frac{g_0 g_1}{\mathrm{FBW}} = 20.63.$$

The SIW width follows from a TE10 cutoff near 0.7 f0 through the Cassivi–Wu correction $a_{\mathrm{siw}} = a_{\mathrm{eff}} + d^2/(0.95p)$ [21]. The cavity lengths follow from the TE101 resonance. Each feed is a 50 Ω microstrip line (Hammerstad [22]) flared into the SIW by a Deslandes taper [23]. At the board edge it becomes a grounded coplanar waveguide (CPWG) launch sized to the SMA footprint: a 1.2 mm trace with 0.2512 mm gaps, which solves to 50.0 Ω by conformal mapping [24], [25]. The coupling-matrix formulation [26], [27] provides both the ideal target response and, later, the diagnostic fit.

## Toolchain and working method

Fig. 1 shows the flow and the division of labour. Fig. 2 shows how the effort was distributed in time.

*Phase 1: analytic design, openEMS and CST*

The first iteration (June 2026) used analytic dimensions, the free FDTD solver openEMS [28] for calibration, and CST Studio Suite 2021 [29] as the reference. Python generated a parametric CST VBA macro that builds native solids wrapped in the solver's history list. The analytic cavity length of 38.23 mm put the CST passband at 2.818 GHz with a 293 MHz band, 3.65 dB insertion loss and only 1.02 dB return loss. The external coupling was uncalibrated and heavily over-coupled.

openEMS was abandoned for three reasons. First, FDTD resolves the 1.5 mm cylindrical vias on a staircase grid, which needs at least three cells per diameter over the whole board. The resonant, high-Q structure also rings down slowly, so each full-filter run took minutes to tens of minutes. Second, the ports did not give a trustworthy S11. The lumped port radiated like a point source instead of launching the microstrip mode. The microstrip port launched the mode correctly, but its S11 de-embedding was unstable and even returned $|S11| > 1$. The passband and S21 were reproduced, but the match could not be tuned, and matching was the whole remaining problem. Third, there was no eigenmode solver for the coupling calibration of Sec. IV, which had to be emulated by reading peak splits off driven sweeps.

Palace [30], an MPI-parallel finite-element solver built on MFEM [31] with gmsh meshes [32], removed all three obstacles. Its tetrahedral mesh conforms to the vias. Its wave ports solve the 2D port mode, as CST's do, and are passive by construction. It has an eigenmode solver, and it is driven entirely by JSON files, which makes it easy for an agent to operate. A first Palace model reproduced the CST result for the same geometry and then replaced openEMS in the loop.

*Phase 2 onward: an agent-operated toolchain*

From August the agent consolidated three copied optimizer scripts into one package with a `Topology` interface and a single command-line tool (`siw`). The tool covers synthesis, layout, eigenmode and driven solves, Bayesian design, 1-D scans, calibration, comparison figures and CAD export. Every production number in this paper is reproducible through that tool, and curated results are stored in a `golden/` directory together with their S-parameters and the command that reproduces them. A project brief file (`CLAUDE.md`) acted as the agent's persistent memory across sessions. Handoff documents carried the state from one machine to another.

The ML layer is a scikit-optimize GP [33], [34] with two fidelities: coarse-mesh evaluations for the search and a fine-mesh confirmation of the winner. It also has a fast physics screen that rejects non-resonant candidates before a full solve, a patience-based early stop, and an objective that went through ten revisions (v9–v18). Later revisions fitted the whole response to the ideal Chebyshev target instead of scoring individual points.

**Compute.** To speed up the search, the work moved to a workstation with an NVIDIA GTX 1660 SUPER (6 GB). Palace ran with its GPU backend on that card, and the author supervised the long agent sessions remotely through Claude Code's Remote Control. The GPU was 1.24× faster than 4 CPU ranks on the dual-mode model, with S-parameters identical to 10  dB. However, the card ran out of memory above about 176 k tetrahedra (174.6 k finished, 178.4 k did not). The 166 mm cascade exceeded that at any adequate mesh density, so no GPU mesh preset was ever validated for it. The GPU was not tried on the fine-mesh verification either, whose CPU run needed 315 k unknowns and about 5.5 GB. All production runs therefore ended up on CPUs with 4 MPI ranks. Four attempts on a 3.7 GB host failed, three of them explicitly out of memory; the final runs used a 16 GB host at about 8 s per frequency point. In short, memory rather than arithmetic was the binding constraint throughout.

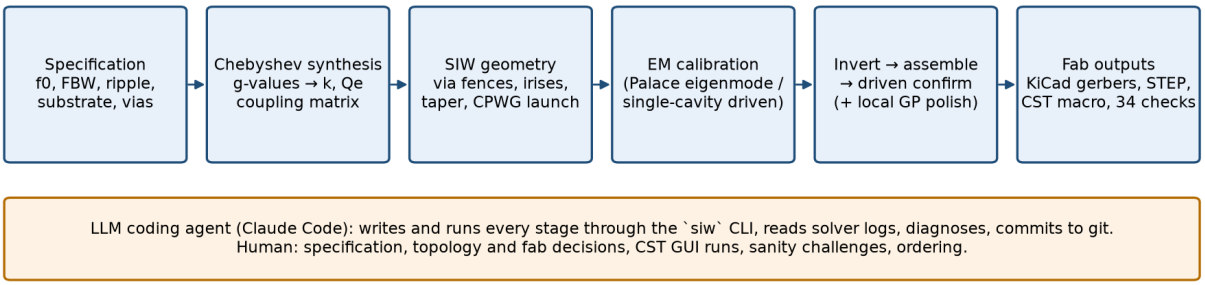



Fig. 1. Design flow. The agent writes and operates every stage through one command-line tool (`siw`); the engineer supplies the specification and the decisions, and runs the commercial-solver cross-checks.

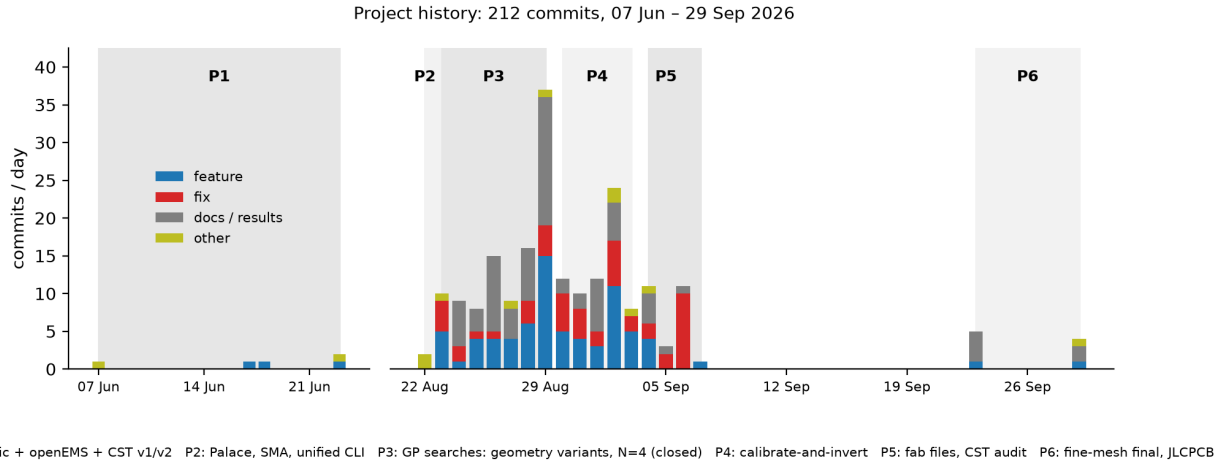



Fig. 2. Commit history by type. P1 was a human-led prototype committed in a few large commits. From P2 onward the agent drove the work, and fixes (red) cluster where new physics or a new tool entered the loop.

## Why global Bayesian search failed

The GP worked as intended in one regime: as a local polish around a validated design. For the dual-mode filter, a 30-evaluation search in a ±2 mm, 5-dimensional box confirmed the landed geometry. Neighbouring probes came back worse, and the apparent 0.41 dB gain lay inside the measured coarse-mesh noise floor of 0.64 dB.

Run globally over the assembled cascade, the same machinery never produced a worst in-band S11 better than −5 dB in more than 60 driven evaluations. The cause is structural. The objective scores the passband width, and that width is discontinuous in the external aperture $w_{\text{ext}}$. Below a coupling threshold only the centre pole clears the −3 dB line, so the band measures one resonance. Above the threshold all N poles merge and the width jumps. On the N = 3 cascade the widths recorded during the search were 16, 22, 36, 65 and then 215 MHz, with the 143.5 MHz target inside the jump. We re-measured this for the paper on the fabricated geometry (Fig. 3c, data in the repository). The width rises 96, 106, 123 MHz for $w_{\text{ext}}$ = 10–12 mm, jumps to 195 MHz at 13 mm, and then *falls* again as the end cavities are over-coupled and the reflection zeros drop from three to one. The target width is crossed twice. The second crossing, at 16 mm, has one reflection zero and −4 dB of return loss; a bandwidth-driven search, or the scan's own inversion (15.2 mm), lands there. A GP assumes a smooth objective, so it inherits both the jump and the false solution. The search-health data showed the same thing: 57–83 % of cascade candidates were screened out, leaving 2–9 informative points per run, against 34–37 % and about 19 informative points on the dual-mode filter, where the search converged.

The replacement measures each coupling on the smallest structure that exhibits it (Fig. 3):

- **Internal coupling.** A closed two-cavity eigenmode solve splits the cavity mode into f1 and f2, and $k = (f_2^2 - f_1^2)/(f_2^2 + f_1^2)$ [20]. This is smooth in $w_{\text{iris}}$.
- **External Q.** One cavity driven symmetrically through two apertures gives the loaded Q, and $Q_e = 2/(1/Q_L - 1/Q_u)$. Here Qu = $1/\tan\delta$ = 270.3, independently matched by the eigenmode Q. This is smooth and monotonic in $w_{\text{ext}}$.

Inverting the fitted curves at the synthesis targets, as-

sembling the filter and making a final driven fc trim converged in about three solves per curve. The re-run calibration for this paper inverts to $w_{\mathrm{iris}} = 13.85$ mm and $w_{\mathrm{ext}} = 13.53$ mm, against 14.42 and 13.96 mm on the fabricated board, whose end cavities were subsequently trimmed. The resulting internal coupling was within 2 % of target. At that checkpoint the N = 4 centre frequency also landed at 2.870 GHz, although the N = 4 design was later abandoned for its return loss (Sec. VI). The GP was kept only as a local polish seeded with the result (`--seed`). The physics that worked is textbook [20], [27]. What the agent contributed was running the calibrations quickly, and diagnosing why the black-box route had failed.

Two further search failures were caused by the parameterization itself rather than by the search. The via fences were first generated by filtering a fixed grid, so the iris width was *quantized*. The whole cascade search box spanned about two distinct geometries, and the GP was tuning a knob that did nothing (Fig. 4). In addition, N $\geq$ 4 had a single shared iris width and could not express the required k = [0.0415, 0.0329, 0.0415].

## Silent defects: a taxonomy

As of commit `b06bf63` (29 Sep 2026) the history holds 212 commits, 117 of them with an agent co-author trailer, and 47 labelled *fix*. Most corrected defects that never raised an error: the solver converged and the plot looked reasonable, but the number was wrong. Table I groups representative cases. Three rules paid for themselves repeatedly. First, if a fix does not change the number, the diagnosis was wrong. Second, after two wrong inferences, stop inferring and run a one-factor ladder. Third, check every result against what physics allows: insertion loss cannot beat a lossless model, the pole count must equal N, and a lossy stack cannot cost 27 dB.

Two cases changed the course of the project. In the first, when a real SMA coax launch replaced the ideal port, the optimized microstrip width of 2.26 mm turned out to be a 41 Ω line. The optimizer had exploited a port that normalizes to whatever it sits on, and the true 50 Ω width is 1.664 mm. The first coax model also used a full-size 0.82 mm pin, which is thicker than the 0.76 mm board, and returned −0.2 dB S11. That was real physics, not a bug (Fig. 5). In the second, an independent audit against CST showed that the −13.37 dB return loss used to choose the cascade over the dual-mode design was a coarse-mesh artefact. The converged value is about −10.7 dB. The decision survived, but the justification had to be rewritten.

## Topologies, lengths and the dual-mode negative result

**Cavity lengths.** The analytic 38.23 mm cavity placed the band at 2.82–2.905 GHz (openEMS/CST). An openEMS sweep to 35.93 mm and 33.64 mm moved it to 2.982 and 3.077 GHz, a slope of about −37 MHz/mm. Palace, which models the loading by the tapers and external apertures that the analytic formula ignores, shortened the cavities to about 34.4 mm. Calibrate-and-invert then gave unequal cavities, 34.24/35.50/34.24 mm. The fine-mesh coupling fit found the end cavities about 17 MHz low, and a −0.3 mm trim (33.94/35.50/33.94 mm) was chosen. A −0.6 mm trim gained only 0.07 dB and cost 0.4 dB of band-edge loss. The measured sensitivity at the end cavities was about 27 MHz/mm, roughly half the 50 MHz/mm rule of thumb.

**Board size.** The design space was also explored for a smaller board (Figs. 6 and 7). The candidates were a single-cavity dual-mode N = 2 filter (119 mm), the N = 3 cascade (166 mm) and an N = 4 cascade (211 mm). Three other single-cavity shapes were tried: corner-cut, circular-cross and rhombic. None came closer to the specification than the dual-mode cavity (best S11 −6 to −9 dB). A compact N = 4 made from two dual-mode cavities was closed after five mechanisms plus a joint GP search all stopped at two of four poles: offset apertures, near-corner coupling, top-copper slots, discrete posts and quad-mode cavities. A four-cavity single-mode cascade works, but it needs 180–206 mm.

**Dual-mode N = 2.** A single square SIW cavity, a = 57.16 mm, carries a degenerate TE102/TE201 pair. A floating internal via row splits the pair and couples it [35], [36]. Under the matched CPWG launch its best response is fc 2.904 GHz, 316.6 MHz bandwidth, −10.76 dB return loss and 1.45 dB insertion loss. Real copper adds 0.32 dB. One-dimensional scans on the assembled filter show that, on this topology, bandwidth and match are *opposed* on every geometry parameter. At $w_{\mathrm{iris}} = 11$ mm the band is 145.6 MHz (4.94 % FBW), but the return loss is −1.75 dB. The pad-only launch used earlier reached −17.7 dB, but that launch was not the one fabricated. N = 2 has only one internal and one external coupling, which is not enough freedom to place two poles and also match them. The cascade has an extra degree of freedom at 1.4× the length, and it won the head-to-head comparison.

## The fabricated design

The chosen geometry is $l_{\mathrm{cav1}} = 33.939$, $l_{\mathrm{cav2}} = 35.503$, $w_{\mathrm{iris}} = 14.416$, $w_{\mathrm{ext}} = 13.96$, $l_{\mathrm{taper}} = 15.738$ and $w_{\mathrm{ms}} = 1.6636$ mm, on a 165.78 × 45.79 mm board. Table II and Fig. 8 summarise the fine-mesh result (2.5/0.45 mm volume mesh, 0.15 mm coax ports, 9.12 mm air, 61 points).

**Why it misses the specification.** A lossy N = 3 coupling-matrix fit to the refined sweep (Fig. 9) puts the external coupling on target (+0.4 %), which exonerates the launch. It finds the internal coupling 4.8 % weak and the end cavities 17 MHz low. The trim halves the detuning, but the worst in-band point sits at the upper band edge, where the equiripple region ends near 2.935 GHz against a target edge of 2.942 GHz. Adding +0.6 mm of $w_{\mathrm{iris}}$ on top of the trim widened the band without helping. Adding +0.4 mm of $w_{\mathrm{ext}}$ on top of both was strongly harmful, and

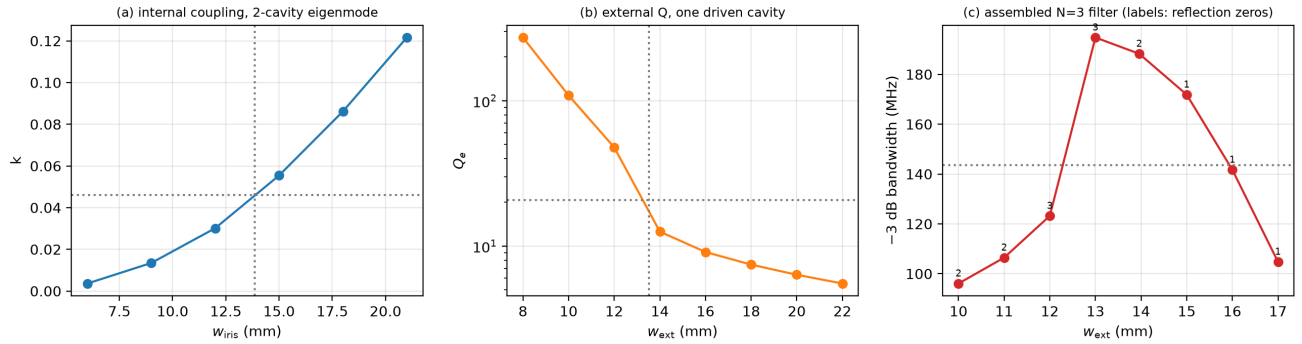



Fig. 3. (a), (b) The calibration observables are monotonic and invertible (dotted: synthesis target and inverted width). (c) The −3 dB bandwidth of the assembled filter, the quantity a black-box objective scores, jumps between 12 and 13 mm and then falls. Labels: number of reflection zeros. Coarse mesh.

TABLE I
REPRESENTATIVE SILENT DEFECTS, HOW THEY WERE CAUGHT, AND THE GUARD ADDED

| Class | Defect (effect) | Caught by | Guard now in code |
|---|---|---|---|
| Parameterization | Iris width quantized by grid filtering (search box ≈2 geometries); single iris width for $N \geq 4$ (spec unreachable); stale via defaults 0.8/1.5 mm | Scan showing identical meshes; coupling targets | Continuous via segments; per-iris parameters |
| Measurement / metric | Bandwidth taken first-to-last crossing, bridging two disjoint lobes; sweep window narrower than the band (fc reported = window centre, 48 MHz error charged nothing); truncated $Q_e$ points giving an inversion outside its own bracket | Identical fc across evaluations; physics sanity | `n_lobes`, `band_truncated`, bracket check |
| Mode selection | Eigenmode pair = smallest gap picked a higher-order cluster: $k = 0.0028$ instead of 0.0583 (21×) | Inspecting the raw eigenvalue list | Window around $f_0$, then the two lowest modes |
| Mesh / boundary tagging | SMA end-cap classified by centroid swallowed the PTFE annulus at $r_{in} \geq 0.3773$ mm (S21 = −79.5 dB); faces left with no boundary condition (IL better than lossless); a domain with no material | Bisection over pin radius; IL > lossless | Face-extent test; unclassified-face report; domain-coverage guard |
| Model validity | Wave port on the trace self-normalizes, so $W_{ms}$ was tuned to a 41 Ω line; surface-impedance BC applied to internal two-sided sheets (−7.8 dB IL on copper); 4.5–20 μm layers meshed as slivers | Real coax launch; Palace warning; one-factor ladder | Real SMA model; exterior-only conductivity; thin layers excluded |
| Plumbing | Adapter dropped soldermask/via-conductivity before the solver (two studies ran with the feature off); N=3 and N=4 results overwrote each other (shared tag); cached curve reused under a retuned label | Result unchanged by a real change | Order in every tag; geometry fingerprint on caches |
| Export | CST macro wrote the launch at 1000× (6 m coax, vias 30 m off-board), twice; taper dropped; coax without shield; PEC end caps on open walls | Cross-check of `.bas` vs mesh | 34 file-level fab checks incl. magnitude test |
| Agent process | VBA properties guessed by name analogy; edits from a parallel agent absorbed unreviewed; comment overclaims; a process declared dead by a flawed `ps` filter | CST errors; adversarial self-review | Only properties seen in a real run; reviewed diffs |
| Convergence | Selection metric (S11 −13.37 dB) from a non-converged mesh; refined model gives −10.7 dB | CST/Palace audit | Fine-mesh final check |

also pulled the end cavities down. No isolated $w_{\text{ext}}$ trial was run. A coordinated local move remains open.

**Why the insertion loss is high.** For a narrowband filter, $\text{IL}_0 \approx 4.343 \sum g_i/(\text{FBW}\, Q_u)$ [20], so the loss is set by the unloaded Q. The Palace runs model the metal as a perfect conductor, so their Q is bounded by the dielectric: $Q_d = 1/\tan\delta = 270$. The coupling-matrix fits returned 229–320, which brackets that bound; the 320 value exceeds it and reflects the fit's uncertainty rather than lower loss. At the simulated 6.3 % bandwidth, $Q_u = 270$ predicts 0.8 dB of dissipation, and the −11 dB in-band mismatch adds 0.35 dB. We attribute the remaining 0.5 dB of the simulated 1.70 dB to radiation from the open microstrip, taper and launch sections, which we did not separate. Real copper lowers Q further. For a TE101 cavity 0.76 mm high, $Q_c \approx 540$ with smooth copper [37], so $Q_u \approx 180$ and the dissipation rises by about 0.4 dB. That is consistent with the 0.59 dB that the exterior-face copper model added. The measured board should therefore show roughly 2.2–2.3 dB. This is close to the floor for this laminate, thickness and bandwidth, not a design error. Each loss mechanism

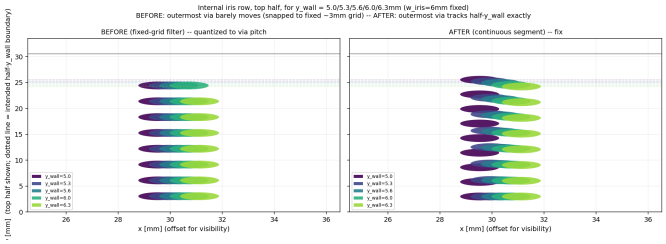



Fig. 4. Via-fence quantization. Filtering a fixed pitch grid collapsed different iris widths onto identical meshes. Segment-wise placement makes the geometry continuous in its parameters.

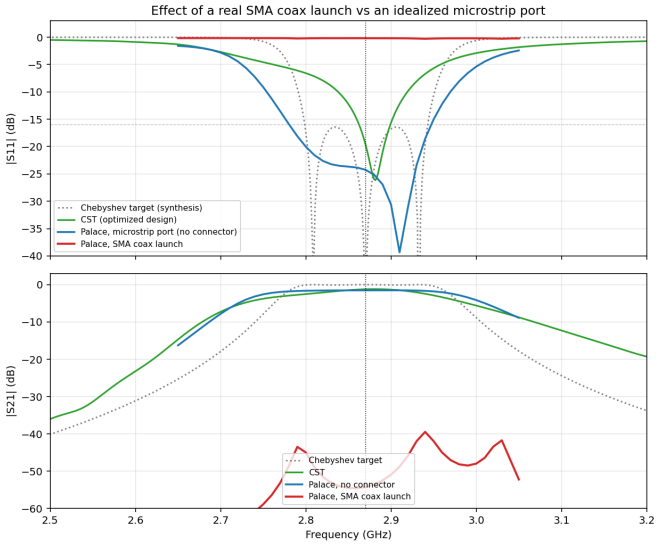



Fig. 5. First real-connector model. An ideal port hid a 41 Ω feed, and a cable-size pin larger than the board shorts the launch.

points to an improvement:

- Qd is fixed by the laminate: $\tan\delta = 0.0037$ is high for a filter.
- Qc scales with the substrate height, and 0.76 mm is thin.
- The loss scales as 1/FBW, so a narrow band costs more.
- The band edges lose more (2.58 dB) because the group delay peaks there.

The 0.59 dB conductor-loss figure comes from a model that applies copper conductivity to the exterior faces only; the internal trace and ground sheets stay ideal, so it is a lower bound. A 1.52 mm RO4350B board only reaches IL ≈ 1.3 dB, because Qd still dominates. RT/duroid 5880 ($\tan\delta \approx 0.0009$) at 1.52 mm would bring it to about 0.5 dB.

**Cross-solver check.** CST Frequency Domain, run from the generated macro, agreed with Palace on fc (2.876 vs. 2.875 GHz) and IL (1.80 vs. 1.85 dB). Getting there took ten debugging steps on the macro and solver setup, seven of them actual fixes, and a coax-only control model to separate port problems from filter problems. The bandwidth agreed less well: 129.6 MHz in CST against 150 MHz in Palace at the reference mesh. The time-domain solver was never made fully stable.

TABLE II
FABRICATED DESIGN: SIMULATED (FINE MESH) VS. SPECIFICATION

| | Target | Simulated |
|---|---|---|
| Centre frequency | 2.870 GHz | 2.876 GHz |
| −3 dB bandwidth | 143.5 MHz | 182.3 MHz |
| Worst in-band $\lvert S_{11}\rvert$ | −16 dB | −11.15 dB |
| Insertion loss, peak / worst in band | — | 1.70 / 2.58 dB |
| Rejection at 2.45 / 3.5 GHz | — | 61 / 39.5 dB |
| Measured | *pending (Sec. VIII)* | |

**The connector is part of the design.** Figs. 10 and 11 compare three launches on the same geometry: a bare microstrip feed, the CPWG launch without a connector, and the CPWG launch with the SMA. Without the connector only one reflection zero remains, and the worst in-band S11 is −6.3 dB (microstrip) or −5.1 dB (CPWG), against −11.2 dB with it. The filter was tuned with the connector's reactance in place. The bare CPWG port also shows a port-modelling trap: its first two Palace modes belong to the ground pours (S21 ≈ −78 dB), and the signal mode is mode 3.

**Manufacturing.** The same command-line tool writes the KiCad board, the gerbers, the drill files, a STEP solid, and 34 automated file-level checks [38]. The checks include drawn-versus-meshed geometry and CST-versus-mesh via positions. JLCPCB stocks RO4350B 0.76 mm but offers only 1 oz outer copper. That leaves the finished thickness at about 0.84 mm at the pads, against the SMA fork slot of 0.81 ± 0.1 mm, which is marginal and can be fixed by scraping 15–20 µm of copper under the connector legs. The no-mask requirement is carried by the gerbers and stated in the order notes. The board has been ordered.

## MEASURED RESULTS

*[Pending — the boards are in manufacture at JLCPCB. This section will report a two-port VNA measurement: |S11|, |S21| and group delay over 2.5–3.3 GHz and 0.1–8 GHz, overlaid on the fine-mesh simulation and the CST result. It will include the measured fc, bandwidth, insertion loss and return loss, the spread across boards, and a re-fit of the coupling matrix to the measurement. Planned de-embedding: SMA-to-SMA through-line if fabricated, otherwise the connector is left in the reference plane, as it is in the simulation.]*

## DISCUSSION: GAPS AND ROOM TO IMPROVE

**What was AI-driven, and what was not.** The agent wrote about 11.5 k lines of design code, launched and monitored the Palace eigenmode and driven solves, read solver logs, localized most of the defects in Table I, and produced the fabrication files. The engineer chose the specification, topology, laminate and vias, decided when to stop tuning, ran every CST check, and more than once rejected a confident but wrong explanation. The design physics that

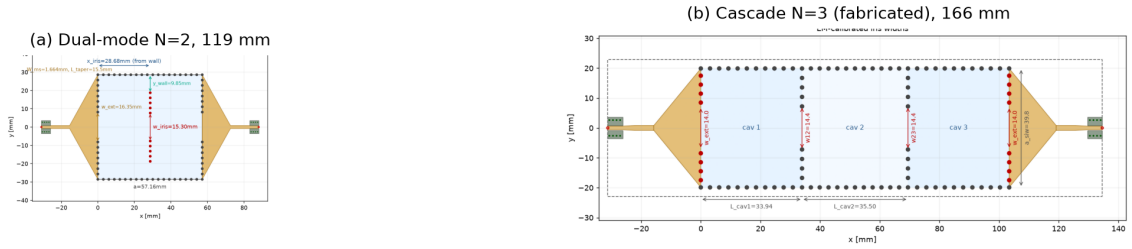



Fig. 6. (a) Dual-mode N=2: one square cavity whose $TE_{102}/TE_{201}$ pair is split by a floating via row. (b) The fabricated N=3 cascade, with external apertures in red and the CPWG/SMA footprints at the board edges.

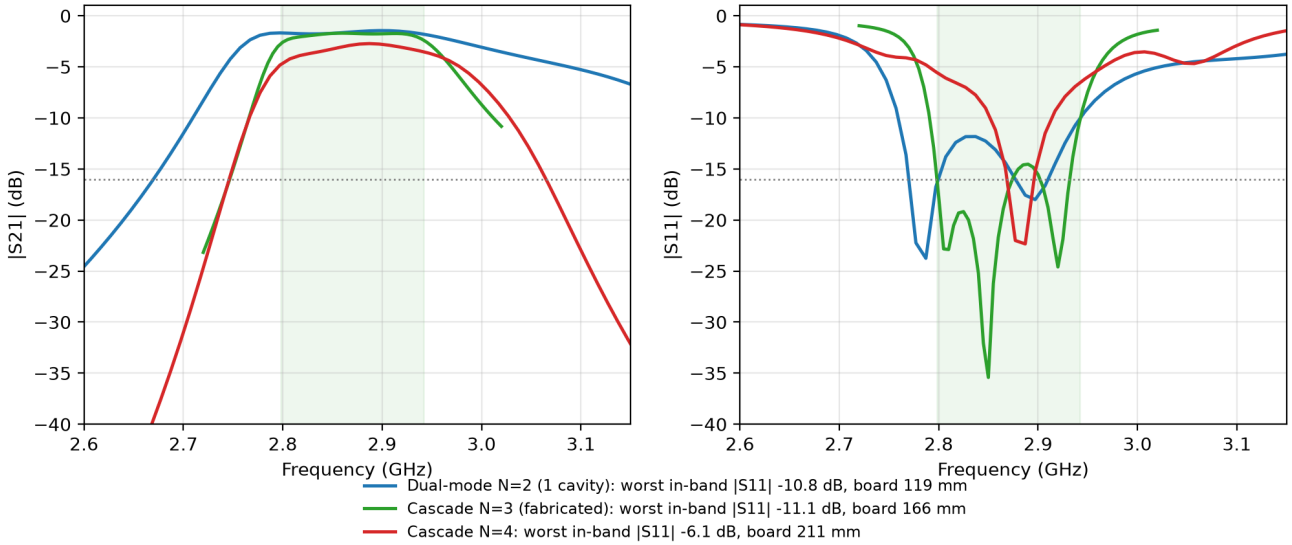



Fig. 7. Best designs of each topology under the same CPWG + SMA launch (Palace). The N=3 curve is the fine-mesh final. The other two are the production mesh.

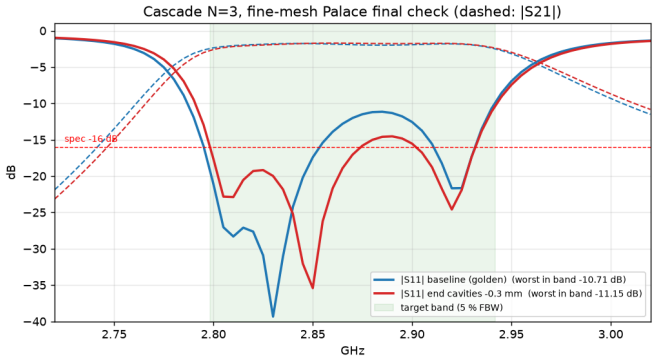



Fig. 8. Fine-mesh final check. The −0.3 mm end-cavity trim lifts the interior ripple from −11.2 to −14.5 dB. The worst point stays at the 2.94 GHz band edge.

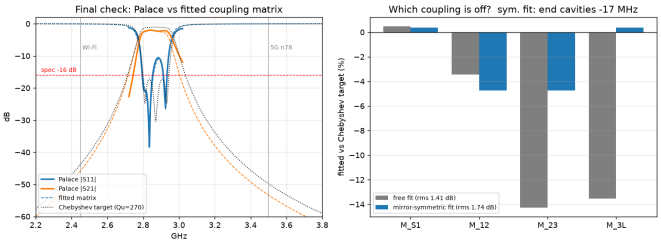



Fig. 9. Lossy coupling-matrix fit used to diagnose the return-loss gap: external Q on target, $k$ 4.8 % weak, end cavities 17 MHz low.

finally worked is classical. The acceleration came from automation and diagnosis, not from new electromagnetics. Stating this plainly is more useful than overclaiming.

**Agent failure modes observed.**

- Confident misdiagnoses: two wrong explanations of one stack failure were each falsified by a single re-run.
- APIs guessed by analogy: VBA properties that do not exist.
- Documentation drift: headline numbers kept in prose after the code had changed. The repository's status files had to be rewritten several times.
- Unreviewed edits from a parallel agent absorbed into a commit.
- Premature optimism about untested geometries.

Each was caught by a physics check, a solver error or a human question, which argues for keeping those checks in the loop.

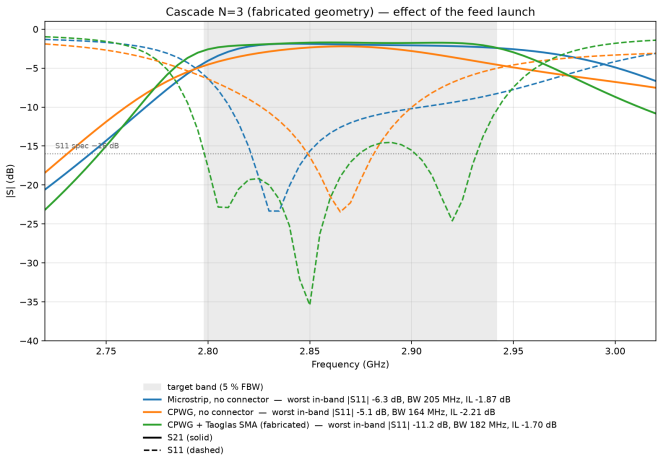



Fig. 10. Same geometry, three launches. Removing the SMA costs ≈5 dB of return loss because the design was tuned with it in place.

**Gaps.**

- No measurement yet.
- The specification is not met: −11.2 dB against −16 dB, and 182 MHz against 143.5 MHz.
- S11 varies by about 5 dB across mesh and port models, while fc, bandwidth and IL are robust. The match is therefore inside the model uncertainty.
- Qe is calibrated with a dielectric-only Qu and is slightly over-estimated.
- Conductor loss is modelled only on exterior faces. Thin fabrication layers (mask, ENIG) could not be resolved.
- A spurious passband at 5.8–6.8 GHz needs a low-pass filter in a real system.
- There is no controlled comparison against an engineer working without the agent, and no accounting of tokens or cost.
- This is a single case study: one engineer, one model family, one filter topology. It shows what can happen, not how often.
- The taxonomy is self-reported by the same engineer–agent pair that made the mistakes. Git corroborates that each fix happened, but not the classification. It also lists only caught defects; the physics checks may still miss a comparable class of error.
- The code has no regression test suite. Its safety net is the physics sanity checks and the automated fabrication checks.

**Improvements.**

- A coordinated local correction of $w_{\text{ext}}$, $l_{\text{cav1}}$ and $w_{\text{iris}}$, driven by the coupling-matrix fit.
- Adaptive frequency sweeps and y-symmetric half-models, for roughly 2× memory and 5–10× fewer solves, which would bring the cascade within GPU limits.
- A GPU with more memory.
- A lower-loss or thicker laminate.
- Space-mapping or surrogate models [2] trained on the calibration data.
- Making the physics sanity checks a mandatory step of the agent loop rather than a habit.

## Conclusion

An LLM coding agent took an SIW bandpass filter from specification to fabrication files, with one engineer supervising. Two things decided the outcome. The first was replacing black-box search over a discontinuous objective with calibration of smooth observables on small structures. The second was catching a long series of silent defects with physics checks. The fabricated N = 3 filter is centred within 6 MHz of target with 1.7 dB loss, close to its laminate's floor. It misses the return-loss and bandwidth specification for reasons the coupling-matrix fit identifies. Measurements will close the loop.

## Acknowledgment

The design software, simulations, figures and a first draft of this text were produced with Claude Code (Anthropic; Claude Sonnet 5, Opus 5 and 5.5, Fable 5) under the author's direction, as described in Sec. III. The author reviewed all results and text and takes full responsibility for them [39]. The complete history and data are available in the project repository.

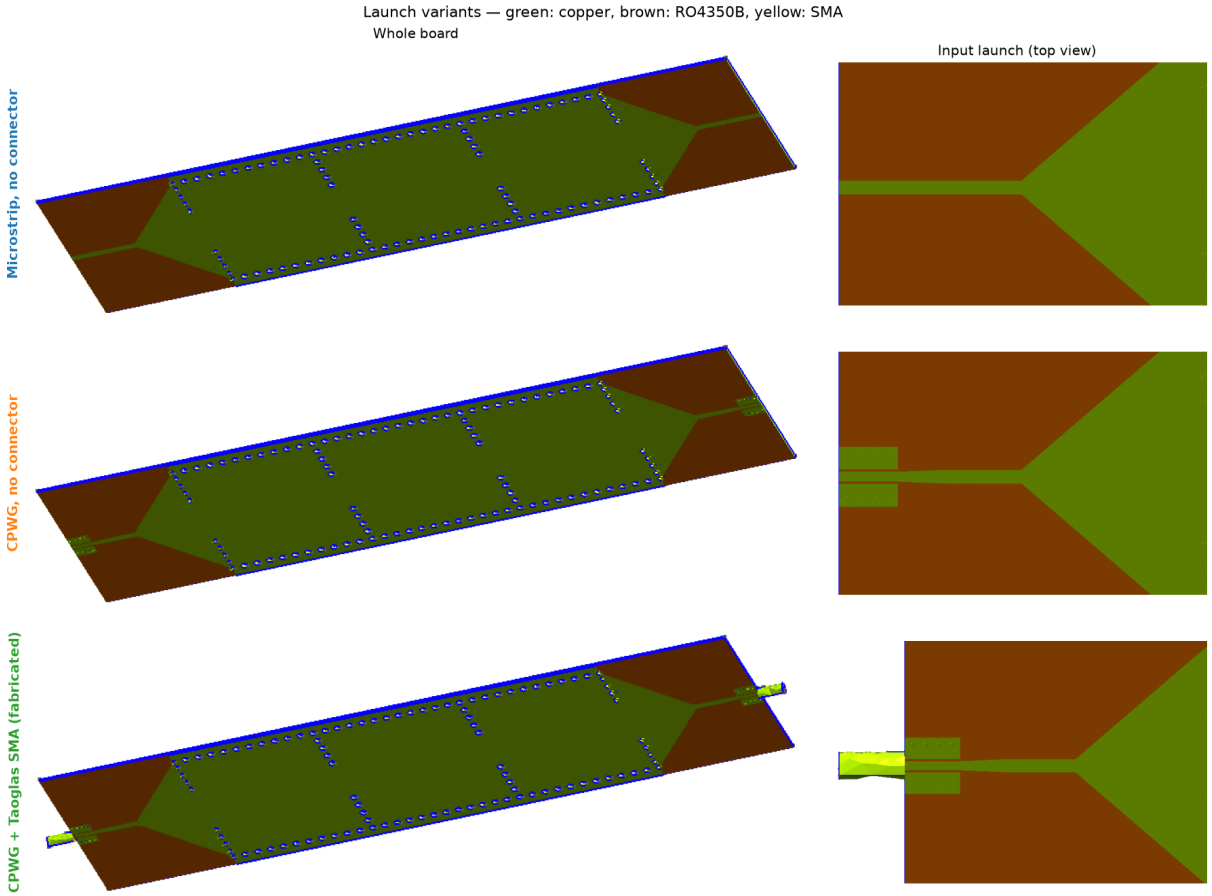



Fig. 11. Palace models of the three launches: whole board and input-launch close-up. Green: copper; brown: RO4350B; yellow: SMA connector.

## References


[1] Q.-J. Zhang and K. C. Gupta, *Neural networks for RF and microwave design*. Norwood, MA: Artech House, 2000.

[2] J. W. Bandler *et al.*, "Space mapping: The state of the art," *IEEE Transactions on Microwave Theory and Techniques*, vol. 52, no. 1, pp. 337–361, 2004, doi: 10.1109/TMTT.2003.820904.

[3] D. R. Jones, M. Schonlau, and W. J. Welch, "Efficient global optimization of expensive black-box functions," *Journal of Global Optimization*, vol. 13, no. 4, pp. 455–492, 1998, doi: 10.1023/A:1008306431147.

[4] J. Snoek, H. Larochelle, and R. P. Adams, "Practical Bayesian optimization of machine learning algorithms," in *Advances in neural information processing systems 25*, 2012, pp. 2951–2959.

[5] T. Sallam and A. M. Attiya, "Fast and accurate extraction of microwave filter coupling matrix via physics-informed deep learning," *Scientific Reports*, vol. 16, p. 17192, 2026, doi: 10.1038/s41598-026-46255-w.

[6] S. Yao *et al.*, "ReAct: Synergizing reasoning and acting in language models," in *International conference on learning representations (ICLR)*, 2023.

[7] J. Yang *et al.*, "SWE-agent: Agent-computer interfaces enable automated software engineering," in *Advances in neural information processing systems 37*, 2024.

[8] Y. Xing *et al.*, "RF-Agent: A practical framework for building language agents for RFIC design." arXiv:2607.18772, 2026.

[9] H. Lu *et al.*, "RFAmpDesigner: A self-evolving multi-agent LLM framework for automated radio frequency amplifier design." arXiv:2605.10093, 2026.

[10] M. Heinrichs, O. Moschner, S. Tewes, V. Wienstroer, A. Sezgin, and R. Kronberger, "From prompt to prototype: Towards a frontier LLM driven RF engineering workflow." arXiv:2608.31006, 2026.

[11] X. Zhang, H. Zhang, and S. H. Tan, "Understanding bugs in modern agentic frameworks: A study of symptoms, root causes, and triggering conditions." arXiv:2604.08906, 2026.

[12] D. Deslandes and K. Wu, "Single-substrate integration technique of planar circuits and waveguide filters," *IEEE Transactions on Microwave Theory and Techniques*, vol. 51, no. 2, pp. 593–596, 2003, doi: 10.1109/TMTT.2002.807820.

[13] M. Bozzi, A. Georgiadis, and K. Wu, "Review of substrate-integrated waveguide circuits and antennas," *IET Microwaves, Antennas & Propagation*, vol. 5, no. 8, pp. 909–920, 2011, doi: 10.1049/iet-map.2010.0463.

[14] X.-P. Chen and K. Wu, "Substrate integrated waveguide filter: Basic design rules and fundamental structure features," *IEEE Microwave Magazine*, vol. 15, no. 5, pp. 108–116, 2014, doi: 10.1109/MMM.2014.2321263.

[15] Anthropic, "Claude Code: Agentic coding tool." ht tps://docs.anthropic.com/en/docs/claude-code, 2026.

[16] Rogers Corporation, "RO4000 series high frequency circuit materials (RO4003C, RO4350B) data sheet." https://www.rogerscorp.com, 2026.

[17] F. Xu and K. Wu, "Guided-wave and leakage characteristics of substrate integrated waveguide," *IEEE Transactions on Microwave Theory and Techniques*, vol. 53, no. 1, pp. 66–73, 2005, doi: 10.1109/TMTT.2004.839303.

[18] D. Deslandes and K. Wu, "Accurate modeling, wave mechanisms, and design considerations of a substrate integrated waveguide," *IEEE Transactions on Microwave Theory and Techniques*, vol. 54, no. 6, pp. 2516–2526, 2006, doi: 10.1109/TMTT.2006.875807.

[19] Taoglas, "EMPCB.SMAFSTJ.B.HT edge-mount SMA female jack, specification SPE-12-8-134/C/WY." Datasheet, 2026.

[20] J.-S. Hong and M. J. Lancaster, *Microstrip filters for RF/microwave applications*. New York: John Wiley & Sons, 2001.

[21] Y. Cassivi, L. Perregrini, P. Arcioni, M. Bressan, K. Wu, and G. Conciauro, "Dispersion characteristics of substrate integrated rectangular waveguide," *IEEE Microwave and Wireless Components Letters*, vol. 12, no. 9, pp. 333–335, 2002, doi: 10.1109/LMWC.2002.803188.

[22] E. Hammerstad and O. Jensen, "Accurate models for microstrip computer-aided design," in *IEEE MTT-s international microwave symposium digest*, 1980, pp. 407–409. doi: 10.1109/MWSYM.1980.1124303.

[23] D. Deslandes, "Design equations for tapered microstrip-to-substrate integrated waveguide transitions," in *IEEE MTT-s international microwave symposium digest*, 2010, pp. 704–707. doi: 10.1109/MWSYM.2010.5517884.

[24] B. C. Wadell, *Transmission line design handbook*. Norwood, MA: Artech House, 1991.

[25] R. N. Simons, *Coplanar waveguide circuits, components, and systems*. New York: Wiley-IEEE Press, 2001.

[26] R. J. Cameron, "Advanced coupling matrix synthesis techniques for microwave filters," *IEEE Transactions on Microwave Theory and Techniques*, vol. 51, no. 1, pp. 1–10, 2003, doi: 10.1109/TMTT.2002.806937.

[27] R. J. Cameron, C. M. Kudsia, and R. R. Mansour, *Microwave filters for communication systems: Fundamentals, design, and applications*, 2nd ed. Hoboken, NJ: John Wiley & Sons, 2018.

[28] T. Liebig, A. Rennings, S. Held, and D. Erni, "openEMS – a free and open source equivalent-circuit (EC) FDTD simulation platform," *International Journal of Numerical Modelling: Electronic Networks, Devices and Fields*, vol. 26, no. 6, pp. 680–696, 2013, doi: 10.1002/jnm.1875.

[29] Dassault Systèmes, "CST Studio Suite 2021." https://www.3ds.com/products/simulia/cst-studio-suite, 2021.

[30] AWS Labs, "Palace: PArallel LArge-scale Computational Electromagnetics." https://github.com/awslabs/palace, 2023.

[31] R. Anderson, J. Andrej, A. Barker, *et al.*, "MFEM: A modular finite element methods library," *Computers & Mathematics with Applications*, vol. 81, pp. 42–74, 2021, doi: 10.1016/j.camwa.2020.06.009.

[32] C. Geuzaine and J.-F. Remacle, "Gmsh: A 3-D finite element mesh generator with built-in pre- and post-processing facilities," *International Journal for Numerical Methods in Engineering*, vol. 79, no. 11, pp. 1309–1331, 2009, doi: 10.1002/nme.2579.

[33] T. Head, M. Kumar, H. Nahrstaedt, G. Louppe, and I. Shcherbatyi, "Scikit-optimize/scikit-optimize." Zenodo, 2021. doi: 10.5281/zenodo.1157319.

[34] C. E. Rasmussen and C. K. I. Williams, *Gaussian processes for machine learning*. Cambridge, MA: MIT Press, 2006.

[35] W. Shen, X.-W. Sun, W.-Y. Yin, J.-F. Mao, and Q.-F. Wei, "A novel single-cavity dual mode substrate integrated waveguide filter with non-resonating node," *IEEE Microwave and Wireless Components Letters*, vol. 19, no. 6, pp. 368–370, 2009, doi: 10.1109/LMWC.2009.2020017.

[36] Z. Xu, G. Zhang, H. Xia, and M. Xu, "Novel hexagonal dual-mode substrate integrated waveguide filter with source-load coupling," *The Scientific World Journal*, vol. 2014, p. 915740, 2014, doi: 10.1155/2014/915740.

[37] D. M. Pozar, *Microwave engineering*, 4th ed. Hoboken, NJ: John Wiley & Sons, 2012.

[38] JLCPCB, "PCB manufacturing capabilities and high-frequency (Rogers) PCB." https://jlcpcb.com/capabilities/pcb-capabilities.

[39] IEEE, "Submission and peer review policies: Guidelines for artificial intelligence (AI)-generated text." https://journals.ieeeauthorcenter.ieee.org, 2023.